



# Introduction to Kubernetes

*Frédéric Lassabe*



UNIVERSITÉ  
FRANCHE-COMTÉ



# What is kubernetes ?



- ▶ Container orchestration system
- ▶ Main features
  - ▶ Service discovery and load balancing
  - ▶ Storage orchestration
  - ▶ Deployments automation
  - ▶ Automated resource allocation
  - ▶ Self-repair
  - ▶ Configuration and secret management
  - ▶ Batch processing
  - ▶ Horizontal scaling
  - ▶ IP/IPv6 support
- ▶ a.k.a. K8s
- ▶ Open source





## Kubernetes overview

The control plane

Workers

Other useful components

## Kubernetes deployment

## Using the cluster



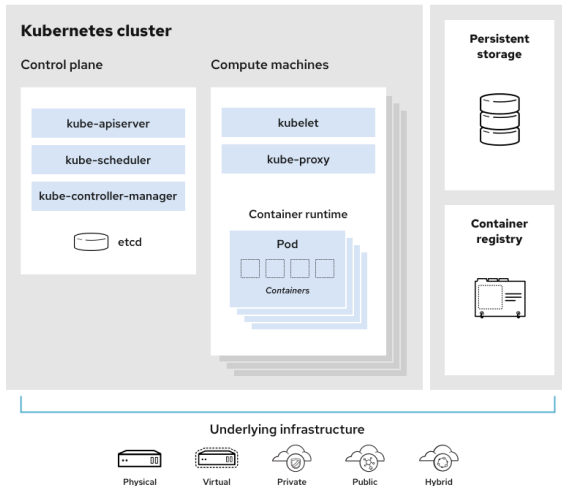
# Cluster



- ▶ A k8s cluster is
  - ▶ a working deployment
  - ▶ a group of hosts
    - ▶ physical
    - ▶ virtual
  - ▶ running linux containers
  - ▶ Two types of hosts :
    - ▶ Control plan
    - ▶ Nodes (workers)



# Architecture



# Workers



- ▶ also named Nodes
- ▶ have their own Linux environment
- ▶ run pods
- ▶ pods made of containers



# Control plane



- ▶ maintaining the cluster state
  - ▶ applications running
  - ▶ container images used
  - ▶ resources allocated
- ▶ takes commands from users (usually admins)
- ▶ orchestrates tasks on nodes



# Cloud-native development



- ▶ Requirement : build and maintain
- ▶ Fast deployment
- ▶ Microservices
  - ▶ Scale better
  - ▶ Easier to modify
  - ▶ Not everything can scale





# Production



- ▶ Production
  - ▶ Applications on multiple containers
  - ▶ Across multiple hosts
- ▶ Kubernetes
  - ▶ Orchestration
  - ▶ Management
  - ▶ To deploy
  - ▶ At scale



# With kubernetes



- ▶ Build application services
- ▶ Spanning multiple containers
- ▶ Scale containers
- ▶ Monitor containers' health
- ▶ Integrates
  - ▶ Network
  - ▶ Storage
  - ▶ Security
  - ▶ Telemetry
  - ▶ And other features




# The need for k8s



- ▶ Scaling application
  - ▶ Multiple containers
  - ▶ Colocated and cooperating
- ▶ Containers
  - ▶ Efficient and self-contained applications modules
- ▶ Microservices
  - ▶ Efficient for app development, deployment, maintenance
  - ▶ Significant increase of numbers of containers
  - ▶ Increased complexity
- ▶ k8s attempts to solve complexity



# Kubernetes features





## Kubernetes overview

The control plane

Workers

Other useful components

## Kubernetes deployment

## Using the cluster



# Role



- ▶ Cluster control
- ▶ Data about cluster state
- ▶ Check applications have sufficient
  - ▶ Containers
  - ▶ Resources
- ▶ Monitors nodes constantly
- ▶ You set an application
- ▶ The control plane guarantees it is executed



# kube-apiserver



- ▶ API to interact with the cluster
- ▶ Control plane's front end
- ▶ Both for requests
  - ▶ internal
  - ▶ external
- ▶ Checks request validity
- ▶ Then executes it
- ▶ Accessible from
  - ▶ REST API
  - ▶ kubectl
  - ▶ kubeadm





- ▶ In charge of
  - ▶ Checking containers health
  - ▶ Finding resources for deploying new containers
- ▶ Uses various metrics
  - ▶ Resources requirements for a pod
  - ▶ Health of the cluster
- ▶ Schedule pods to appropriate node





# kube-controller-manager



- ▶ Controllers run the cluster
- ▶ Controller manager contains several controller features :
  - ▶ Check running pods against scheduler
    - ▶ Can restart failed containers
  - ▶ Route requests to endpoints
  - ▶ Create accounts
  - ▶ Manage API access tokens
  - ▶ etc.





- ▶ As etc manages your linux host configuration
- ▶ etcd is your configuration repository
  - ▶ Key-value database
  - ▶ Fault tolerant
  - ▶ Distributed





## Kubernetes overview

The control plane

**Workers**

Other useful components

## Kubernetes deployment

## Using the cluster



# What are nodes ?



- ▶ Compute resource
- ▶ Hosts services via containers inside pods
- ▶ Pods run on nodes
- ▶ Scheduled from the control plane
- ▶ Scaling up :
  - ▶ Add more nodes
  - ▶ Join them to the cluster



# Pod



- ▶ Smallest unit in kubernetes object model
- ▶ One pod = one instance of application
- ▶ One pod
  - ▶ One container
  - ▶ Or tightly coupled containers
  - ▶ Options for running the containers
- ▶ Pods may have persistent storage



# Container runtime engine



- ▶ k8s manages containers
- ▶ Requires a container runtime
  - ▶ docker (with cri-dockerd)
  - ▶ CRI-O
  - ▶ Containerd
- ▶ You shall choose a runtime for any k8s instance
- ▶ Used runtime shall support the Open Container Interface standard



# kubelet



- ▶ Each node has a kubelet
- ▶ Application to communicate with the control plane
- ▶ Ensures that containers run in pods
- ▶ Instructions from control plane executed by kubelet



# kube-proxy



- ▶ Required in all nodes
- ▶ Network proxy
- ▶ Supports Kubernetes networking services
- ▶ Handles network communications
  - ▶ Internal as well as external
  - ▶ relying either on
    - ▶ OS' packet filtering layer
    - ▶ Directly forwarded by kube-proxy







## Kubernetes overview

The control plane

Workers

Other useful components

## Kubernetes deployment

## Using the cluster



# Persistent storage



- ▶ Containers run applications
- ▶ May need persistent data
- ▶ Stateful applications
- ▶ Hardware storage mapping on nodes would tie a pod to a node
- ▶ Use of network storage
- ▶ Transparent for the pod and its container
- ▶ Persistent network storage outlives pods



# Container registry



- ▶ A storage for container images used by kubernetes
- ▶ Images for your applications
- ▶ Registry owned by the cluster's administrators
- ▶ Or a third party registry



# Underlying infrastructure

- ▶ Cloud infrastructures are not untangible : real hardware
- ▶ Kubernetes deploys either on :
  - ▶ Bare metal
  - ▶ Virtual machines
  - ▶ Public cloud providers
  - ▶ Hybrid cloud environments
- ▶ one or more hosts for *control plan*, with all its services
  - ▶ Redundancy in case of failure
  - ▶ Desirable : at least two control plan nodes
- ▶ one or more different hosts for workers
  - ▶ No colocation of control plan and worker
  - ▶ Workers count depends on :
    - ▶ available resources
    - ▶ Target load





Kubernetes overview

Kubernetes deployment

Control plane

Nodes

Using the cluster



# Constraints



- ▶ Deployments based on
  - ▶ 3 hosts
    - ▶ 1 control plane
    - ▶ 2 nodes
  - ▶ Ubuntu hosts
  - ▶ In VM
- ▶ Based on various documents, mainly
  - ▶ kubernetes website
  - ▶ devopscube.com
  - ▶ Red Hat website





Kubernetes overview

Kubernetes deployment

Control plane

Nodes

Using the cluster



# Steps



- ▶ Install a container runtime
- ▶ Install kubeadm, kubelet, kubectl
- ▶ Initialize
  - ▶ Control plane configuration
  - ▶ Generate the join command for nodes
- ▶ Install a network plugin







Kubernetes overview

Kubernetes deployment

Control plane

Nodes

Using the cluster



# Steps



- ▶ Join the worker to the master node
- ▶ Validate all components (check everything is working)



# And then ?



- ▶ Install a metrics server
- ▶ Use the cluster





Kubernetes overview

Kubernetes deployment

Using the cluster



# Pods



- ▶ Pods are the service-level unit
- ▶ They contain one or more containers
- ▶ Everything in a pod
  - ▶ Tightly coupled
  - ▶ Described by manifests
  - ▶ Orchestrated by kubernetes
- ▶ What goes in a pod or another depends on
  - ▶ Persistence requirements
  - ▶ Application architecture



# Best way to understand k8s



- ▶ Practice
- ▶ Details for the steps are available
  - ▶ In online documentation
  - ▶ In the lab instructions
- ▶ So technical/know-how details aren't in this doc





# Thanks !

